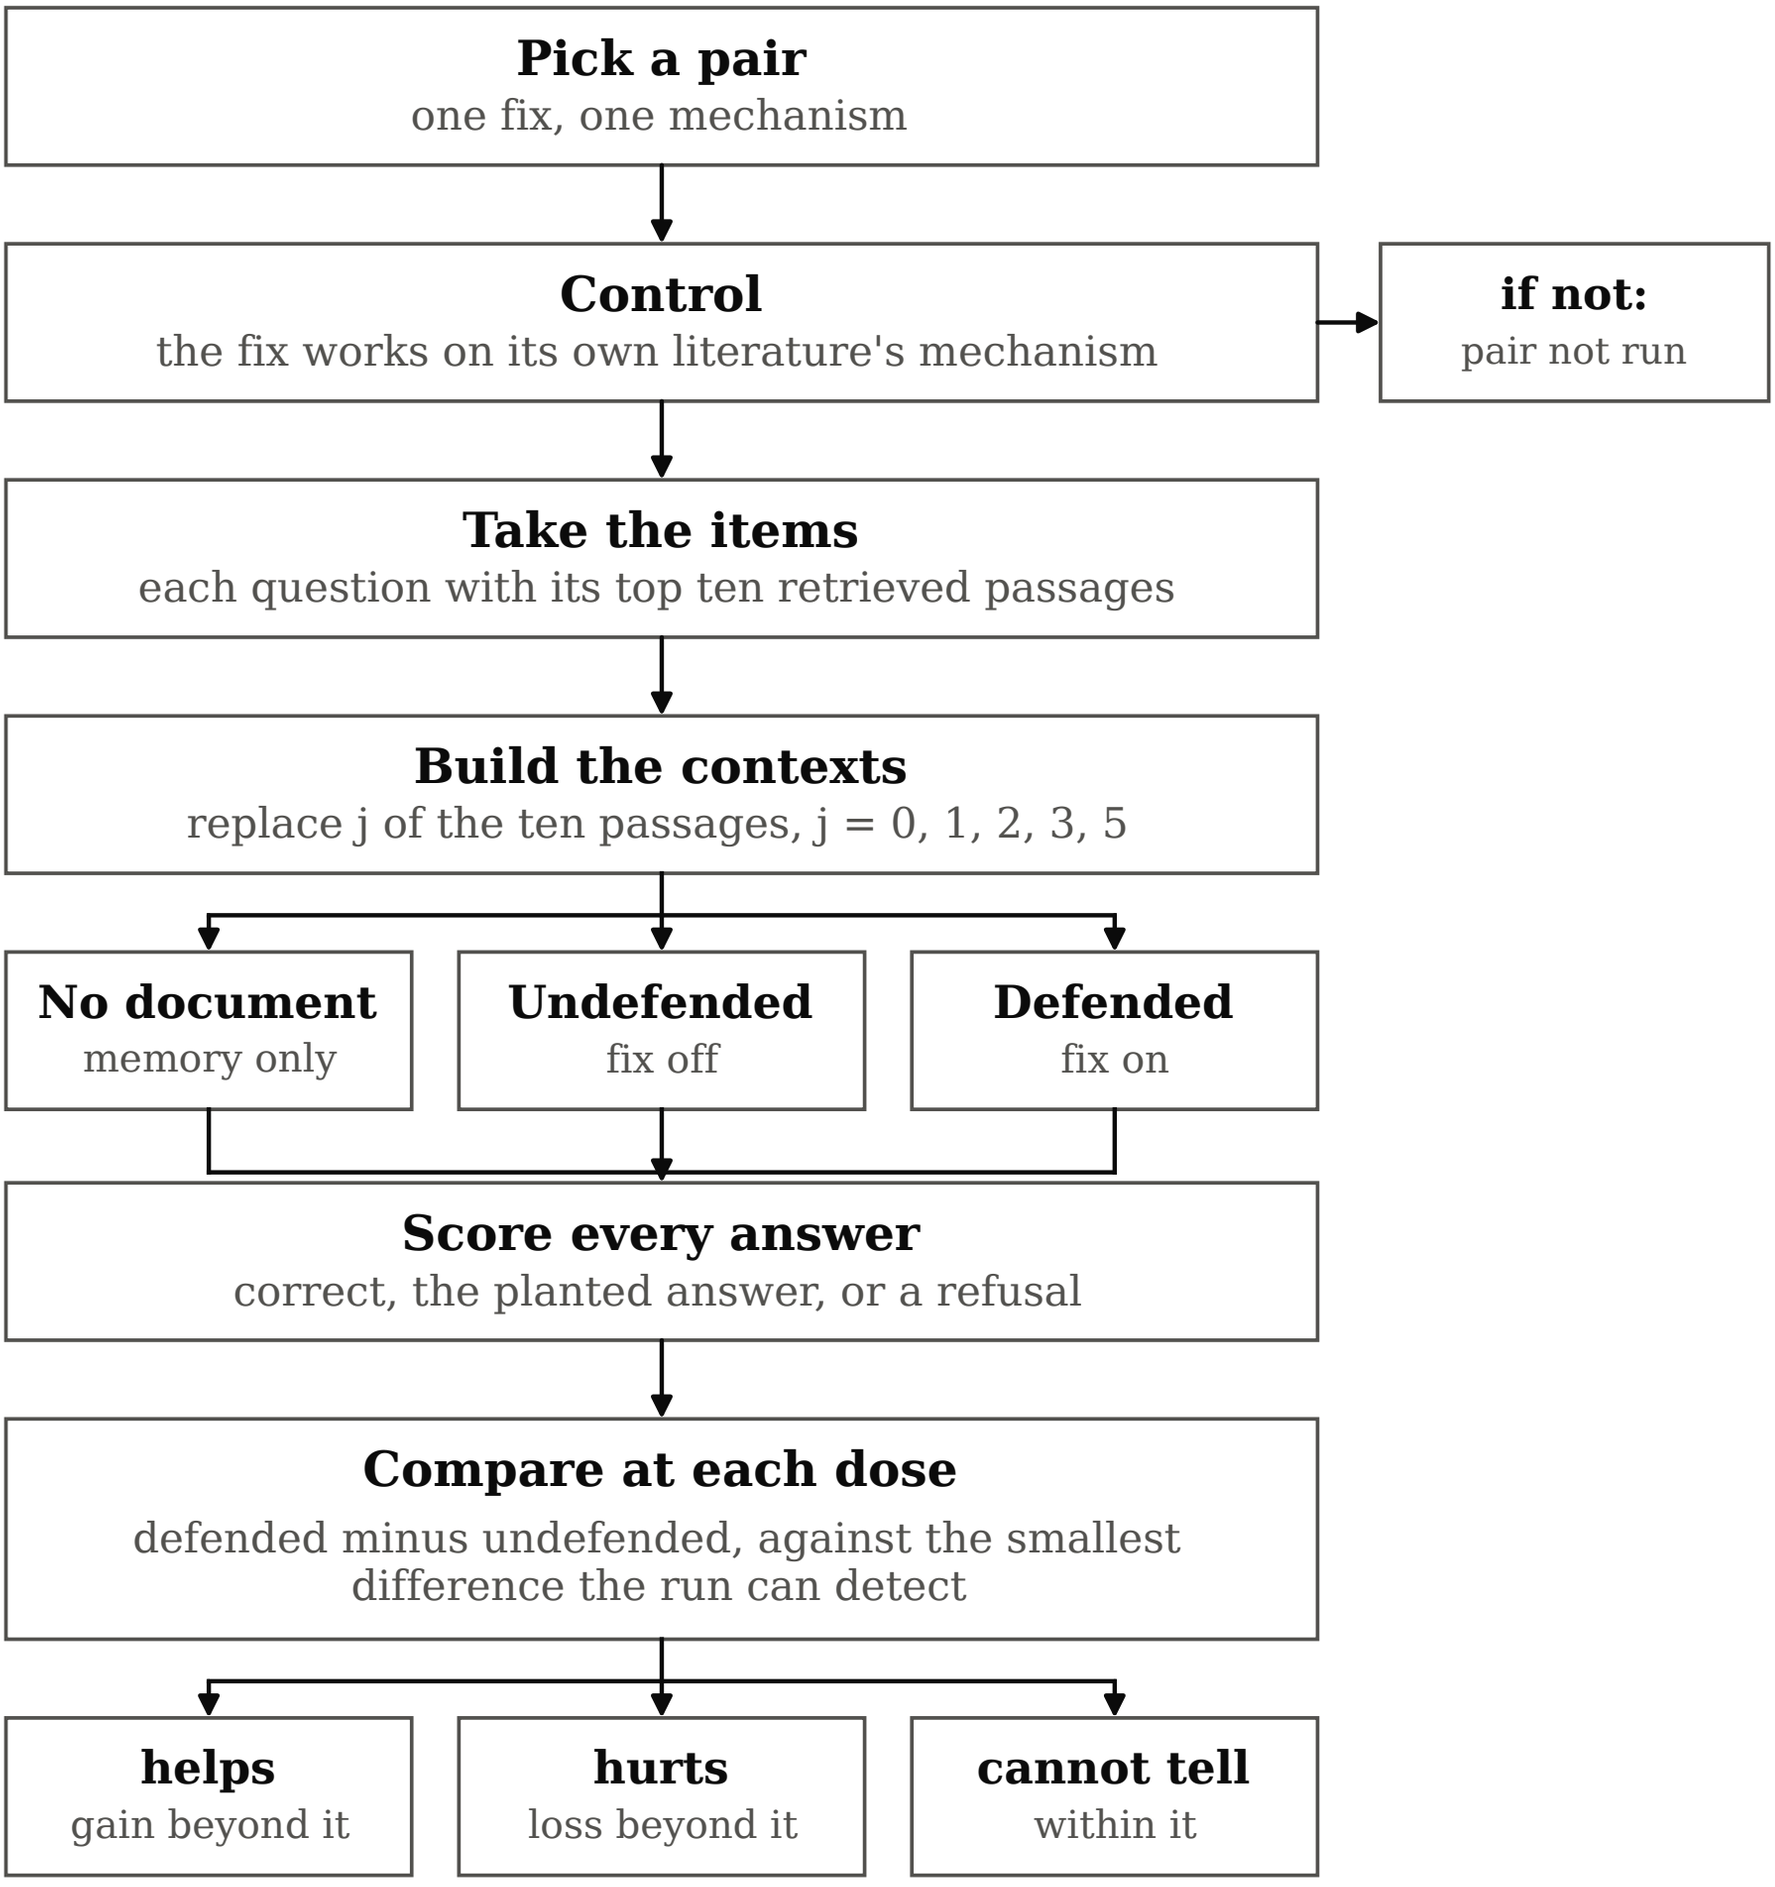
Pick a pair
one fix, one mechanism
Control
the fix works on its own literature's mechanism
if not:
pair not run
Take the items
each question with its top ten retrieved passages
Build the contexts
replace j of the ten passages, j = 0, 1, 2, 3, 5
No document
memory only
Undefended
fix off
Defended
fix on
Score every answer
correct, the planted answer, or a refusal
Compare at each dose
defended minus undefended, against the smallest difference the run can detect
helps
gain beyond it
hurts
loss beyond it
cannot tell
within it